

# Process Creation

## References

1. 'Computer Systems A Programmer's Perspective', Randal Bryant and David O'Hallaron, Pearson Education
2. Unix System Programming, Keith Haviland, Dina Gray and Ben Salama, Addison-Wesley

# Introduction

- Unix Provides a library routine called `system` which allows a shell command to be executed from within a program.
- One can invoke a standard shell as an intermediary, rather than attempt to run the command directly
- - `c` argument used in the invocation of the shell tells it to take commands from the next string argument, rather than standard input.

# Introduction

- Define PATH properly.
- Execute any shell script, you have developed as under:

```
bash-2.03$ cc -o execcommand execcommand.c
bash-2.03$ execcommand
```

```
Enter command to be executed: filecopy
```

```
Enter the filename(s) to be copied
```

```
execcommand.c
```

```
Enter the destination for file(s) to be copied :
```

```
execcommand.bak
```

```
command execution completed
```

```
$ ls -la
```

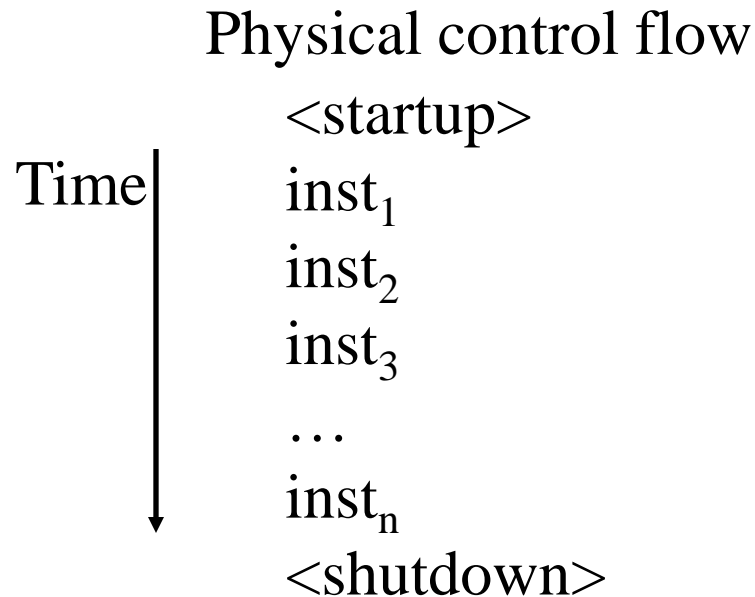
Shell creates a process to run this command

```
$ ls -la | more
```

- Two processes will be created to run this command
- Several processes can run concurrently run the same program

# Control Flow

- Computers do Only One Thing
  - From startup to shutdown, a CPU simply reads and executes (interprets) a sequence of instructions, one at a time.
  - This sequence is the system's physical *control flow* (or *flow of control*).



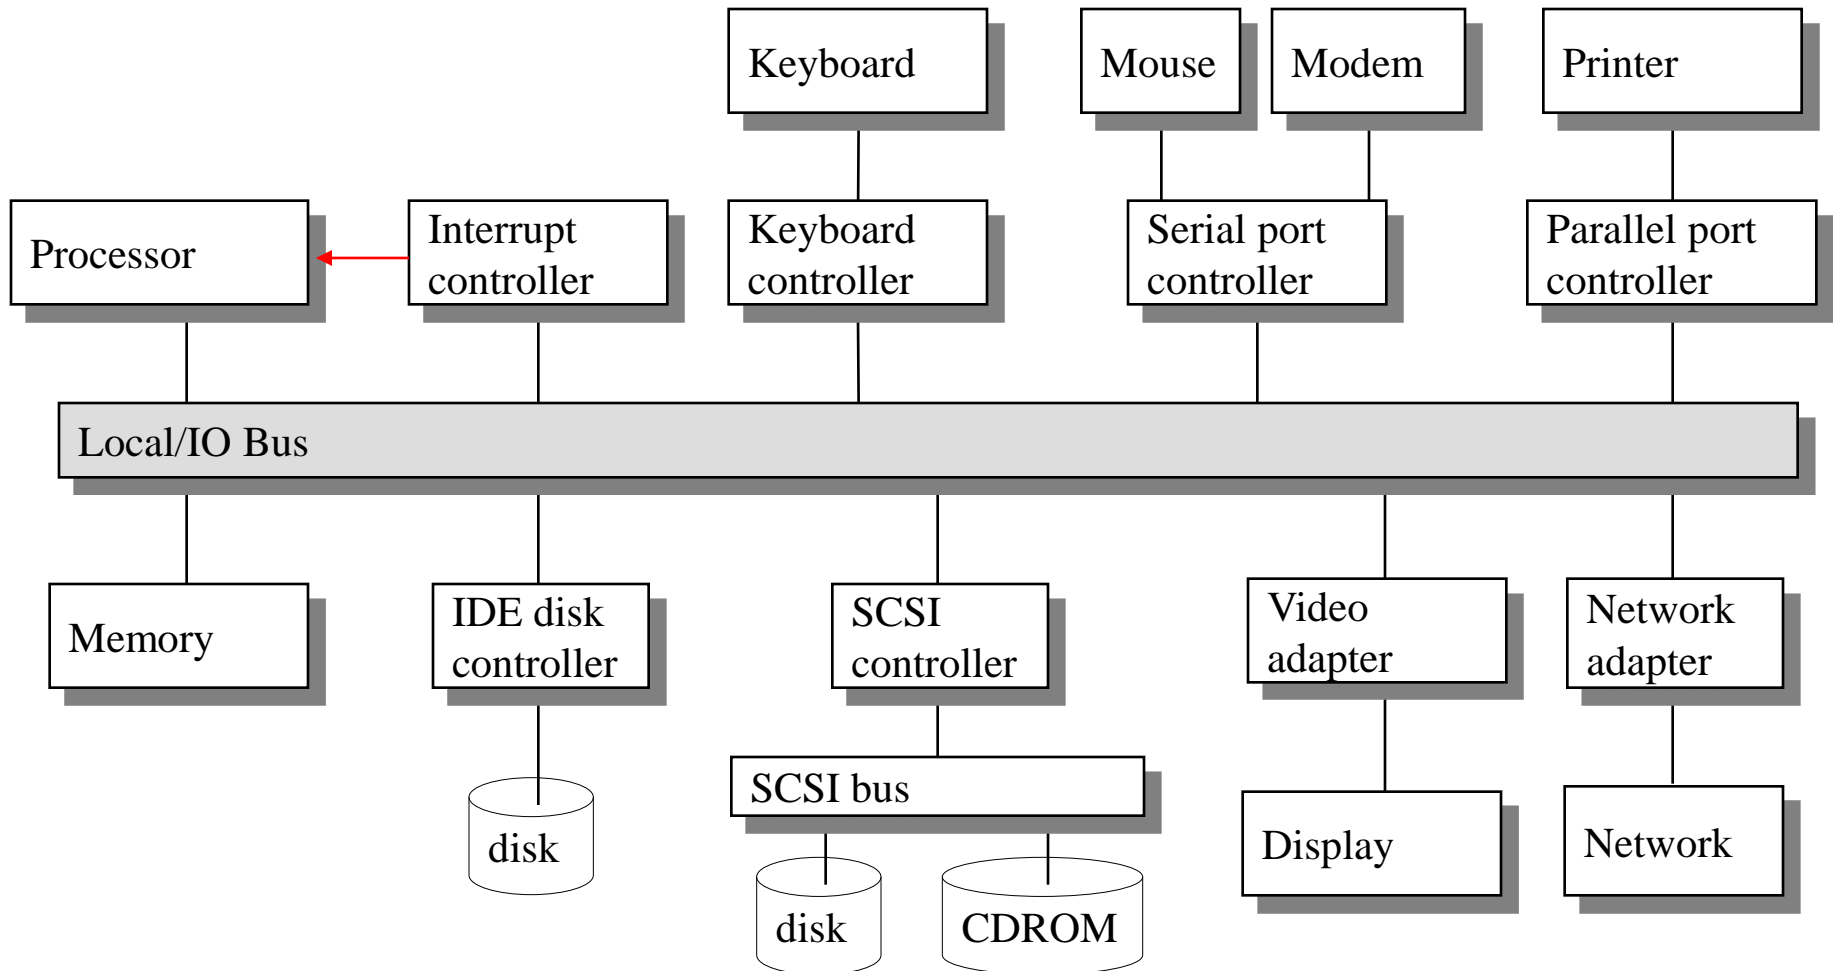
# Altering the Control Flow

- Up to Now: two mechanisms for changing control flow:
  - Jumps and branches
  - Call and return using the stack discipline.
  - Both react to changes in program state.
- Insufficient for a useful system
  - Difficult for the CPU to react to changes in system state.
    - data arrives from a disk or a network adapter.
    - Instruction divides by zero
    - User hits ctrl-c at the keyboard
    - System timer expires
- System needs mechanisms for “exceptional control flow”

# Exceptional Control Flow

- Mechanisms for exceptional control flow exists at all levels of a computer system.
- Low level Mechanism
  - exceptions
    - change in control flow in response to a system event (i.e., change in system state)
  - Combination of hardware and OS software
- Higher Level Mechanisms
  - Process context switch
  - Signals
  - Nonlocal jumps (setjmp/longjmp)
  - Implemented by either:
    - OS software (context switch and signals).
    - C language runtime library: nonlocal jumps.

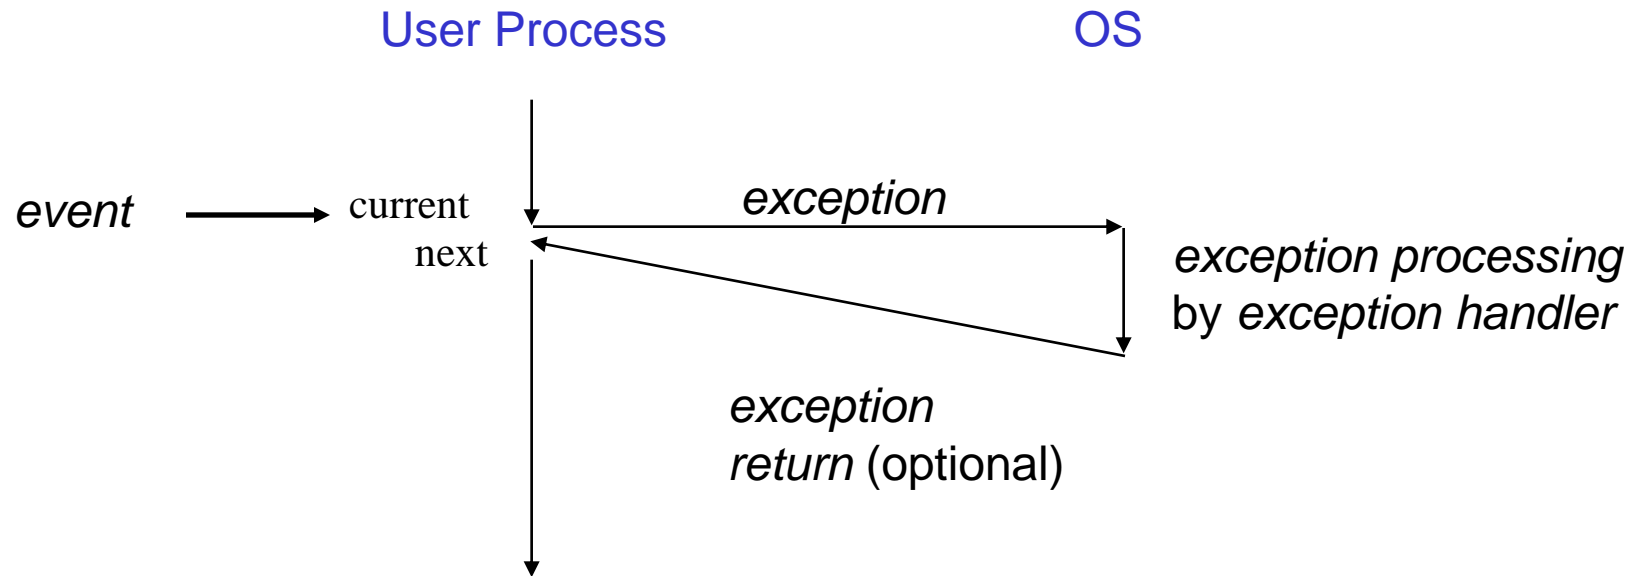
# System context for exceptions



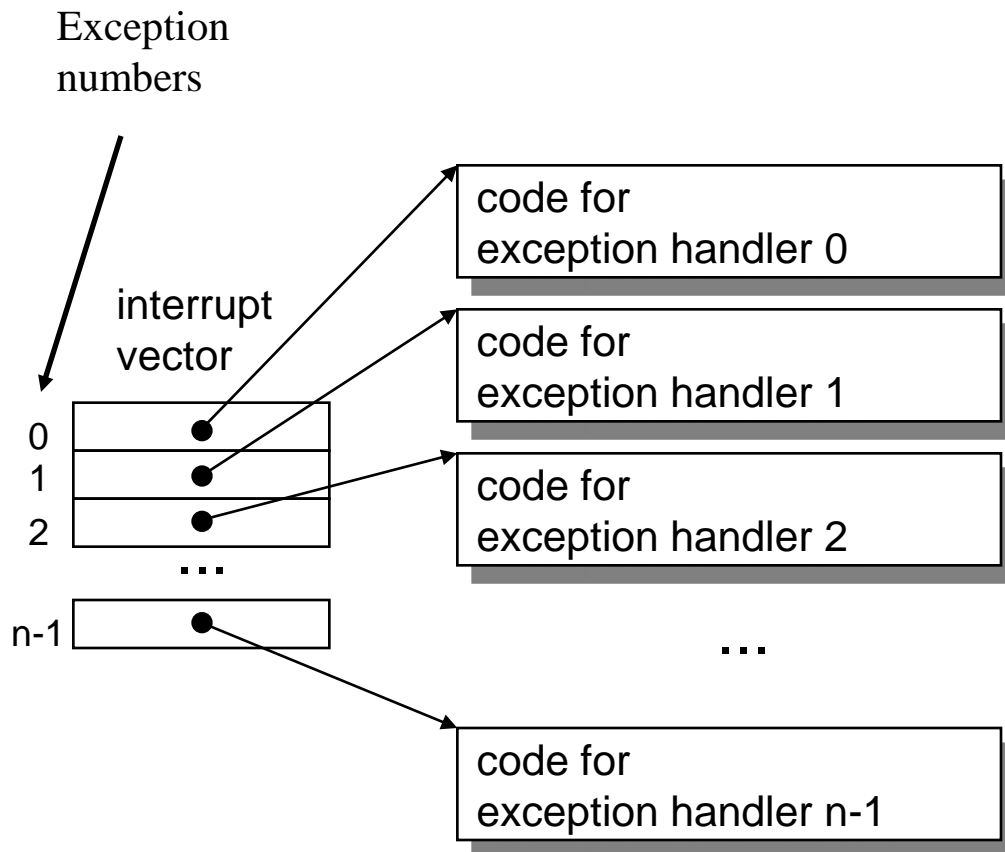


# Exceptions

- An *exception* is a transfer of control to the OS in response to some *event* (i.e., change in processor state)



# Interrupt Vectors



- Each type of event has a unique exception number  $k$
- Index into jump table (a.k.a., interrupt vector)
- Jump table entry  $k$  points to a function (exception handler).
- Handler  $k$  is called each time exception  $k$  occurs.

# Asynchronous Exceptions (Interrupts)

- Caused by events external to the processor
  - Indicated by setting the processor's interrupt pin
  - handler returns to “next” instruction.
- Examples:
  - I/O interrupts
    - hitting ctl-c at the keyboard
    - arrival of a packet from a network
    - arrival of a data sector from a disk
  - Hard reset interrupt
    - hitting the reset button
  - Soft reset interrupt
    - hitting ctl-alt-delete on a PC

# Synchronous Exceptions

- Caused by events that occur as a result of executing an instruction:
  - Traps
    - Intentional
    - Examples: system calls, breakpoint traps, special instructions
    - Returns control to “next” instruction
  - Faults
    - Unintentional but possibly recoverable
    - Examples: page faults (recoverable), protection faults (unrecoverable).
    - Either re-executes faulting (“current”) instruction or aborts.
  - Aborts
    - unintentional and unrecoverable
    - Examples: parity error, machine check.
    - Aborts current program


# Trap Example

- Opening a File

- User calls `open(filename, options)`

```
0804d070 <__libc_open>:  
. . .  
804d082:      cd 80                int     $0x80  
804d084:      5b                   pop     %ebx  
. . .
```

- Function `open` executes system call instruction `int`
    - OS must find or create file, get it ready for reading or writing
    - Returns integer file descriptor



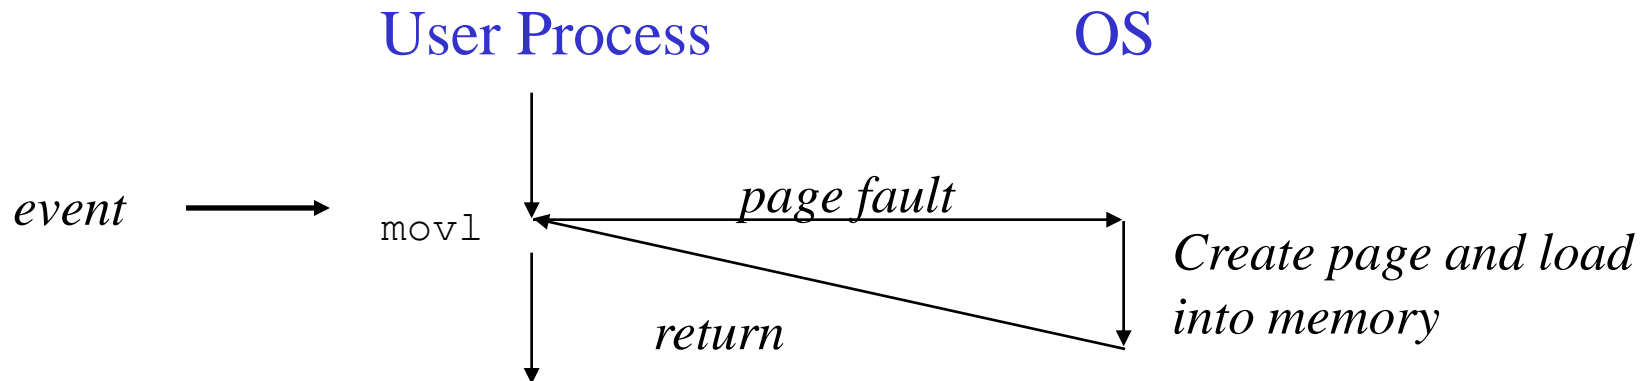
# Fault Example #1

- Memory Reference
  - User writes to memory location
  - That portion (page) of user's memory is currently on disk

```
int a[1000];  
main ()  
{  
    a[500] = 13;  
}
```

```
80483b7: c7 05 10 9d 04 08 0d movl $0xd,0x8049d10
```

- Page handler must load page into physical memory
- Returns to faulting instruction
- Successful on second try



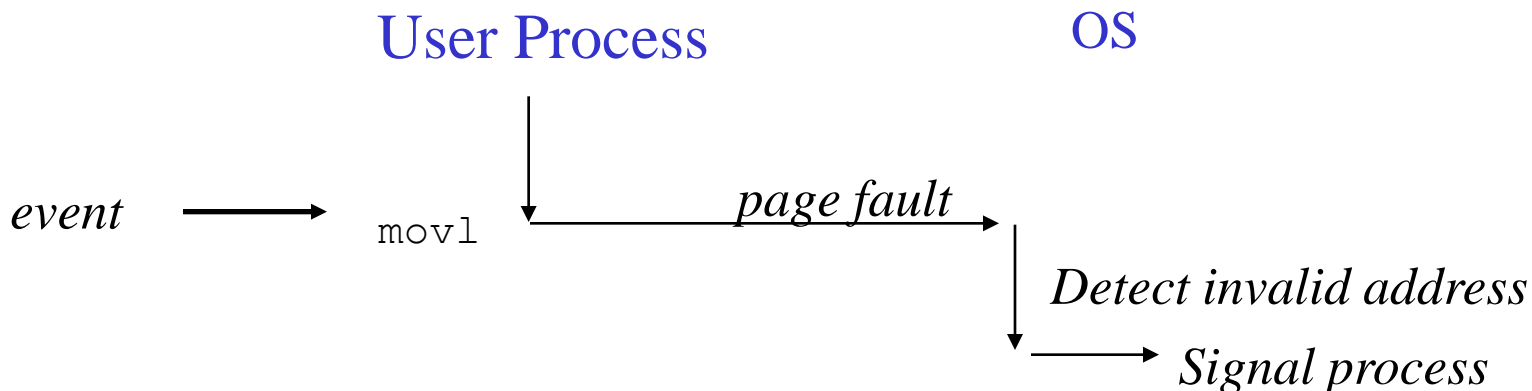
# Fault Example #2

```
int a[1000];  
main ()  
{  
    a[5000] = 13;  
}
```

- Memory Reference
  - User writes to memory location
  - Address is not valid

```
80483b7:      c7 05 60 e3 04 08 0d  movl    $0xd,0x804e360
```

- Page handler detects invalid address
- Sends SIGSEGV signal to user process
- User process exits with “segmentation fault”



# **System Calls for Process Creation and Manipulation**

- **fork:** to create a new process by duplicating the calling process. The basic process creation primitive
- **exec:** A family of library routines and one system call, each of which performs the same underlying function: the transformation of a process by overlaying its memory space with a new program



# System Calls for Process Creation and Manipulation

- wait: provides rudimentary process synchronization. It allows one process to wait until another related process finishes.
- exit: to terminate a process

# fork system call

- Mechanism to transform unix into a multitasking system
- It causes the kernel to create new process called a “child process”
- A child process created with fork is an almost perfect copy of its parent.
- The child process will retain the values they held in the parent, except the return value from fork itself.
- Data available to the child occupies a different absolute place in memory, thus subsequent changes in one process will not affect the variables in the other.

# fork system call

- After fork, the parent process and child process will continue to execute in parallel/concurrently.
- Both will resume execution at the statement immediately after the call to fork.

```
#include <sys/types.h>
#include <unistd.h> /* for fork() */
main()
{
    pid_t pid;          /*holds process-id in parent*/
    printf("One\n");
    pid=fork();
    printf("Two\n");
}
```

```
[sanjay@dslabsrv17 unixprgs]$ fork01
```

One

Two

Two

- Before fork, process A is existing
- After fork process A and Process B will exist, B is the new process spawned by the call to fork and two processes will exist

# fork system call

- `fork` is called without arguments and returns `pid_t`, which is an integer
- The value of `pid` distinguishes parent and child.
- In parent, `pid` is set to non-zero, positive number.
- In child it is set to zero.
- Return value in parent and child differs, the programmer is able to specify different actions for the processes.
- The number returned to the parent in `pid` is called process-id of the child

# fork system call

- Both the processes will run concurrently and without synchronization.
- Fork is useful when parent and child perform different but related tasks, cooperating by using one of the unix inter-process communication mechanisms such as signals or pipes.

```
[sanjay@dslabsrv17 unixprgs]$ cc -o runforkexec1  
runforkexec1.c
```

```
[sanjay@dslabsrv17 unixprgs]$ runforkexec1
```

# Inherited data and file descriptors

- All the files open in the parent process are also open in the child process.
- The child will maintain its own copy of the file descriptors associated with each file.
- However, files kept open across a call to fork remain intimately or closely connected in child and parent.
- This is because the read-write pointer for each file is maintained by the system, it is not embedded explicitly within the process itself.

# Inherited data and file descriptors

- When a child process moves forward in a file, the parent process will also find itself at the new position

Question:

- What will happen within a parent process when a child process closes a file descriptor inherited across a fork?



# exec and open files


- Open file descriptors are also normally passed across calls to exec.
- Files open in the original program are kept open when an entirely new program is started through exec.
- The read-write pointers for such files are unchanged by the exec call.
- **fcntl** function can be used to set the close-on-exec flag associated with a file.

# Running new program with exec

- To initiate the execution of a new program
- All varieties of exec performs the same function:
  - Transform the calling process by loading a new program into its memory space
- If exec is successful, the calling program is completely overlaid by the new program, executed from its beginning
- The new process retains process-id of the calling process
- Exec does not create a new subprocess to run concurrently with the calling program
- There is no return from a successful call to exec

# exec family tree

The difference is in the way parameters can be passed to these functions



The real system call

# exec system call

- There are various forms of the exec system call:
- The execl forms use a null-terminated list of arguments;
  - execl() needs a full pathname for the program to execute
  - execlp() makes use of the user's search path so only needs a filename of the executable
  - The execv() forms need an array of arguments rather than a list;

# exec system call

- As with the `execl`, there is an `execp()` from that uses the search path.
- `exec` calls overwrite the existing process so there is no return value from a successful call.
- The arguments to `exec` form the arguments to `main()` of the executed program.

# Use of exec and fork

- By forking and then using exec within the newly created child, a program can run another program within a subprocess and without obliterating or destroying itself
- Refer following programs:
  - exec01.c
  - execlfail.c
  - execlfail.c

```
#include <sys/types.h>
#include <sys/wait.h> /* for wait() */
#include <unistd.h> /* for fork() */
#include <stdio.h> /* for printf() */
#include <stdlib.h> /* for perror() */
```

```
int main(void)
{
    int fatal(char *);
    pid_t pid;
```

```

switch(pid = fork())
{ case -1:
    fatal("fork failed");
    break;
case 0:
    /* child process calls exec */
    execl("/bin/ls", "ls", "-l", (char *)0);
    fatal("exec failed");
    break;
default:
    /* parent process uses wait to suspend execution
    * until child process finishes */
    wait((int *)0);
    printf("ls completed\n");
    exit(0);
}
}
int fatal(char *s)
{    perror(s);
    exit(1);    }

```



# Terminating processes with the exit system call

- Exit system call is used to terminate a process.
- A process will also stop when it runs out of program by reaching the end of the main function, or when main executes a return statement.
- The single, integer argument to exit is called the process' exit status, the low-order eight bits of which are available to the parent process, providing it has executed a wait system call.
- The value returned through exit in this way is normally used to indicate the success or failure of the task performed by the process.

# Terminating processes with the `exit` system call

- By convention a process returns zero on normal termination, some non-zero value if something has gone wrong.

`exit()` has a number of other consequences:

- All open file descriptors are closed.
- If the parent process has executed a wait call, it will be restarted.
- Exit will also call any programmer-defined exit handling routines and perform what are generally described as clean-up actions, e.g. be concerned with buffering in the standard I/O library.
- A programmer can also set at least 32 bit handling routines with the `atexit` function.

# Synchronizing processes

The `wait()` system call

- `wait` temporarily suspends the execution of a process while a child process is running.
- Once the child process is finished, the waiting parent is restarted.
- If more than one child is running then `wait` returns as soon as any one of the parent's offspring exits.
- `wait` is often used by a parent process just after a call to `fork`, e.g.
- Refer: `runforkexecl.c`

# Synchronizing processes

- The combination of fork and wait is most useful when the child process is intended to run a completely different program by calling exec.
- The return value from wait is normally the process-id of the exiting child.
- If wait returns (pid\_t)-1, it can mean that no child exists and in this case errno will contain the error code ECHILD.
- The parent can sit in a loop waiting for each of its offspring.
- When the parent realizes that all the children have terminated, it can continue.

```
#include <sys/types.h>
#include <sys/wait.h> /* for wait() */
#include <unistd.h> /* for fork() */
#include <stdio.h> /* for printf() */
#include <stdlib.h> /* for perror() */
```

```
int main(void)
{
    int fatal(char *);
    pid_t pid;
```

```
    printf("PID before fork, ie PID of current
process: %d\n", getpid());
    switch(pid = fork())
    {
    case -1:
        fatal("fork failed");
        break;
```

**Parent and Child  
Process IDs known  
to each other**

```
case 0:
    /* child process calls exec */
    printf("Message from a Child Process\n");
    printf("\n");
    printf("The Value PID assigned by pid-t and
known to a Child Process : %d\n", pid);
    printf("The value of PID of child process:
%d\n", getpid());
    printf("The value of Parent PID, who created me:
%d\n", getppid());

    execl("/bin/ls", "ls", "-l", (char *)0);
    printf("\n");

    fatal("exec failed");
    break;
```

```
default:
    /* parent process uses wait to suspend execution
    * until child process finishes */
    wait((int *)0);
    printf("Message from a parent process\n");
    printf("PID of child known to the parent: %d\n", pid);
    printf("TASK ACCOMPLISHED, ie. We are back from MOON
\n");
    exit(0);
}
```

```
[sanjay@dslabsrv17 unixprgs]$ cc -o runforkexeccltest runforkexeccltest.c
```

```
[sanjay@dslabsrv17 unixprgs]$ runforkexeccltest
```

PID before fork, ie PID of current process: **24097**

Message from a Child Process

The Value PID assigned by pid-t and known to a Child Process : **0**

The value of PID of child process: **24098**

The value of Parent PID, who created me: **24097**

Output of ls command will be displayed here

Message from a parent process

PID of child known to the parent: **24098**

TASK ACCOMPLISHED, ie. We are back from MOON

# Synchronizing processes

- `wait` takes one argument `status`, a pointer to an integer.
- If the pointer is `NULL` then the argument is simply ignored.
- If `wait` is passed a valid pointer, `status` will contain useful status information when `wait` returns. This information will be the exit-status of the child passed through `exit`.
- Refer: `status01.c`



# Synchronizing processes

- The value returned to the parent via `exit` is stored in the high-order eight bits of the integer status.
- To be meaningful the low-order bits must be zero
- `WIFEXITED` macro defined in `<sys/wait.h>` tests to see if this is in fact the case. It returns the value stored in the high-order bits of status
- If it returns 0 then the child process was stopped in its track by another process using IPC mechanism called **signal**

# Waiting for a particular child process: `waitpid`

- The `exit` system call is used to terminate a process.
- A process also stop when it runs out of program by reaching the end of the `main` function, or when `main` executes a `return` statement.
- The single, integer argument to `exit` is called the process' exit status, the low-order eight bits of which are available to the parent process, providing it has executed a `wait` system call.

# Waiting for a particular child process: `waitpid`

- The value returned through `exit` in this way is normally used to indicate the success or failure of the task performed by the process.
- By convention, a process returns zero on normal termination.
- Some non-zero value indicates something has gone wrong.
- `exit` will close all open file descriptors, if the parent process has executed a `wait` call, it will be restarted.

# Waiting for a particular child process: waitpid

- A process waits for its children to terminate or stop by calling the waitpid function.

```
#include <sys/types.h>
```

```
#include <sys/wait.h>
```

```
pid_t waitpid(pid_t pid, int *status, int  
options);
```

- First argument, pid specifies the process-id of the child process that the parent wishes to wait for.
- If it is set to -1 and the options argument is set to 0, then waitpid behaves exactly the way wait behaves.
- -1 indicates an interest in any child process

# Waiting for a particular child process: waitpid

- If pid is greater than 0 then the parent waits for the child with a process-id of pid.
- Status will hold the status of the child process when waitpid returns.
- The final argument, options, can take a variety of values defined in <sys/wait.h>. Options can take a variety of values defined in <sys/wait.h>. WNOCHANG is the most useful
- It allows waitpid to sit in a loop monitoring a situation but not blocking if the child process is still running.
- If WNOCHANG is set then waitpid will return 0 if the child has not yet terminated

Waiting for a particular child process:  
waitpid

- Refer: waitpid1.c, status03.c

# Waiting for a particular child process: waitpid

- The default behaviour can be modified by setting options to various combinations of the WNOHANG and WUNTRACED.
- WNOHANG: return immediately (with a return value of 0) if none of the child processes in the wait set has terminated yet.
- WUNTRACED: Suspend execution of the calling process until a process in the wait set becomes terminated or stopped. Returns the PID of the terminated or stopped child that caused the return.

# Waiting for a particular child process: `waitpid`

- `WNOHANG` | `WUNTRACED`: Suspend execution of the calling process until a child in the wait set terminates or stops, and then return the PID of the stopped or terminated child that caused the return. Also, return immediately (with a return value of 0) if none of the processes in the wait set is terminated or stopped.



# Checking the Exit Status of a Reaped Child

- If the status argument is non-null, then `waitpid` encodes status information about the child that caused the return in the status argument. The `wait.h` include file defines several macros for interpreting the status argument:
- `WIFEXITED(status)`: Returns true if the child terminated normally, via a call to `exit` or a `return`.
- `WEXITSTATUS(status)`: Returns exit status of a normally terminated child. This status is only defined if `WIFEXITED` returned true.

# Checking the Exit Status of a Reaped Child

- `WIFSIGNALED(status)`: Returns true if the child process terminated because of a signal that was not caught.
- `WTERMSIG(status)`: Returns the number of signal that caused the child process to terminate. This status is only defined if `WIFSIGNALED(status)` returned true.
- `WIFSTOPPED(status)`: Returns true if the child that caused the return is currently stopped.

# Checking the Exit Status of a Reaped Child

- `WSTOPSIG(status)`: Returns the number of the signal that caused the child to stop. This status is only defined if `WIFSTOPPED(status)` returns true.
- Error Conditions:
- If the calling process has no children, then `waitpid` returns -1 and sets `errno` to `ECHILD`. If the `waitpid` function was interrupted by a signal, then it returns -1 and set `errno` to `EINTR`.

# Zombie processes and Premature exits

- A child exits when its parent is not currently executing wait
- A parent exits when one or more children are still running
- A zombie process occupies a slot in a table maintained by the kernel for process control, but does not use any other kernel resources

# Process Attributes

- **Process-id**

```
pid=getpid();
```

```
ppid=getppid();
```

# Process groups and process group-ids

- Unix allows processes to be placed into groups

```
who | awk '{print $1}' | sort -u
```

- It is useful when a set of processes are doing inter-process communication (IPC) with signals
- Each process group is denoted by a process group-id of type `pid_t`, i.e. `pid_t getpgrp(void)`;
- Use `getprgrp()` system call to obtain process group-id.

## Changing process group

- A process can be placed in a new process group

```
int setpgid(pid_t pid, pid_t pgid);
```

# Sessions and session Ids

- Each process belongs to a session
- A session is a collection of a single foreground process group using the terminal and one or more background process groups

```
pid_t getsid(pid_t pid) ;
```

```
pid_t setsid(void) ;
```

- If is passed a value of 0 then it returns the session-id of the calling process otherwise session-id of the process identified by pid is returned.
- Useful for daemons, as they do not have controlling terminal. It can start a sessions and move into a new session.

# The environment

- One can get the environment of a process by adding an extra parameter `envp` to the parameter list of the main function within a program
- Example: `showenv.c`
- The default environment is passed to a process that was created through a call to `exec` or `fork`.



# Tools to manipulate processes

- Unix system provides a number of useful tools for monitoring and manipulating processes:
- `strace`: prints a trace of each system call invoked by a program and its children. You need to compile your program with `-static` to get a cleaner trace without a lot of output related to shared library
- `ps`: Lists processes (including zombies) currently in the system
- `top`: Prints information about the resource usage of current processes
- `kill`: Sends a signal to a process. Useful for debugging programs with signal handlers
- `/proc`: A virtual file system that exports contents of numerous kernel data structures in an ASCII form that can be read by user programs.

# Tools to manipulate processes (cont)

```
[sanjay@dslab66 tmp]$ cc -static -o sigtalk sigtalk.c
[sanjay@dslab66 tmp]$ strace -p 28259
write(1, "\nPARENT: sending SIGQUIT\n\n", 26) = 26
rt_sigprocmask(SIG_BLOCK, [CHLD], [], 8) = 0
rt_sigaction(SIGCHLD, NULL, {SIG_DFL}, 8) = 0
rt_sigprocmask(SIG_SETMASK, [], NULL, 8) = 0
nanosleep({10, 0}, {10, 0})           = 0
munmap(0x400000000, 4096)              = 0
exit_group(0)                         = ?
```

# To see current load average on your linux system

```
[sanjay@dslab66 tmp]$ cat /proc/loadavg  
1.56 0.58 0.64 3/58 28264
```